Back to AI Programming with Python Nanodegree

Image Classifier Application

REVIEW | CODE REVIEW 4 | HISTORY

▼ train.py 2

```
1 import argparse
2
```

SUGGESTION

The code is really well done. You can extract some of the helper functions in both train/predict scripts anc module.
This could further improve readability.

```
3 import torch
4 from torch import nn
5 from torch import optim
6 from torch.autograd import Variable
7 from torchvision import datasets, transforms, models
8 from collections import OrderedDict
9
10
11 def parse_args():
12     """ Parses commandline arguments and assigns default values if needed"""
13     parser = argparse.ArgumentParser(description="Ezhil's Image Classifier")
14     parser.add_argument('data_dir', action="store")
15     parser.add_argument('--save_dir', action='store',
16                         default='checkpoint.pth', dest='checkpoint_dir')
17     supported_arch = ['densenet121', 'vgg13', 'vgg16']
18     parser.add_argument('--arch', action='store',
19                         default='densenet121', choices=supported_arch, dest='ar
20     parser.add_argument('--learning_rate',
21                         action='store', default='0.001',
22                         dest='learning_rate', type=float)
23     parser.add_argument('--hidden_units',
24                         nargs='+', type=int, default=[], dest='hidden_units')
25     parser.add_argument('--epochs',
26                         action='store', dest='epochs', default=3, type=int)
27     parser.add_argument('--gpu',
28                         action='store_true', dest='gpu')
29
30     args = parser.parse_args()
31     if len(args.hidden_units) == 0:
32         args.hidden_units = [400]
33
34     print()
35     print("Starting to train using these inputs")
36     print("{0: <30}".format('Data Directory:'), args.data_dir)
37     print("{0: <30}".format('Checkpoint Dir:'), args.checkpoint_dir)
38     print("{0: <30}".format('arch:'), args.arch)
39     print("{0: <30}".format('learning_rate:'), args.learning_rate)
40     print("{0: <30}".format('hidden_units:'), args.hidden_units)
41     print("{0: <30}".format('epochs:'), args.epochs)
42     print("{0: <30}".format('gpu:'), args.gpu)
43     print()
44     print()
45     return args
46
47
48 def load_data(data_dir):
49     """Loads Data from the training and validation sets and do necessary transf
50     train_dir = data_dir + '/train'
51     valid_dir = data_dir + '/valid'
52
53     train_data_transforms = transforms.Compose([transforms.RandomRotation(30),
54                                                 transforms.RandomResizedCrop(22
55                                                 transforms.RandomHorizontalFlip
56                                                 transforms.ToTensor(),
57                                                 transforms.Normalize([0.485, 0.
58                                                                       [0.229, 0.
59
```

```
    valid_data_transforms = transforms.Compose([transforms.Resize(225),
                                                transforms.CenterCrop(224),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.
                                                                      [0.229, 0.

    train_image_data = datasets.ImageFolder(train_dir, transform=train_data_tra
    valid_image_data = datasets.ImageFolder(valid_dir, transform=valid_data_tra

    return train_image_data, valid_image_data


def get_loader(train_image_data, valid_image_data):
    """Creates training and validation DataLoaders"""
    train_loader = torch.utils.data.DataLoader(train_image_data, batch_size=64,
    valid_loader = torch.utils.data.DataLoader(valid_image_data, batch_size=32)

    return train_loader, valid_loader


def create_classifier(hidden_units):
  layers = []
  total = len(hidden_units)
  for idx, features in enumerate(hidden_units):
    if idx+1 == total:
      layers.append(('output', nn.LogSoftmax(dim=1)))
    else:
      name = 'fc'+str(idx+1)
      layers.append((name, nn.Linear(features, hidden_units[idx+1])))
      if idx+2 < total:
        relu_name = 'relu'+str(idx+1)
        dropout_name = 'dropout'+str(idx+1)
        layers.append((relu_name, nn.ReLU()))
        layers.append((dropout_name, nn.Dropout(p=0.5)))
  print()
  print('Using the classifier as below')
  print(layers)
  print()
  return nn.Sequential(OrderedDict(layers))


def create_model(arch, hidden_units):
    """Creates model from the architecture and hidden units provided in the inp
    model = getattr(models, arch)(pretrained=True)
    for param in model.parameters():
        param.requires_grad = False
    hidden_units.insert(0, model.classifier.in_features)
    hidden_units.append(102)
    classifier = create_classifier(hidden_units)
    model.classifier = classifier
    return model


def train(model, learning_rate, epochs, gpu):
    """Trains the model and prints the running
        and validation losses, validation accuracy"""
    criterion = nn.NLLLoss()
    optimizer = optim.Adam(model.classifier.parameters(), lr=learning_rate)
    if gpu:
```

SUGGESTION

Along with the user provided GPU check, you should also have a condition to check if GPU is available so a

```
        model.cuda()

    print_every = 40
    running_loss = 0
    step = 0

    for e in range(epochs):
        model.train()
        for ii, (inputs, labels) in enumerate(train_loader):
            step += 1
            optimizer.zero_grad()
            inputs, labels = Variable(inputs), Variable(labels)
            if gpu:
                inputs, labels = inputs.cuda(), labels.cuda()

            outputs = model.forward(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.data[0]

            if step % print_every == 0:
                model.eval()
                accuracy = 0
                valid_loss = 0
                for idx, (v_inputs, v_labels) in enumerate(valid_loader):
                    v_inputs, v_labels = \
                        Variable(v_inputs, volatile=True), Variable(v_labels,
                    if gpu:
                        v_inputs, v_labels = v_inputs.cuda(), v_labels.cuda()
```

Rate this review

```
                    v_outputs = model.forward(v_inputs)
                    valid_loss += criterion(v_outputs, v_labels).data[0]
                    ps = torch.exp(v_outputs).data
                    equality = (v_labels.data == ps.max(1)[1])
                    accuracy += equality.type_as(torch.FloatTensor()).mean()
                print("Epoch: {}/{}.. ".format(e + 1, epochs),
                      "Training Loss: {:.3f}.. ".format(running_loss / print_ev
                      "Validation Loss: {:.3f}.. ".format(valid_loss / len(vali
                      "Validation Accuracy: {:.3f}".format(accuracy / len(valic
                running_loss = 0
                model.train()
    print('Model Training Done')
    return model


def save_model(model, train_image_data, args):
    """Saves the mode in the checkpoint directory file mentioned"""
    model.class_to_idx = train_image_data.class_to_idx
    checkpoint = {'hidden_units': args.hidden_units,
                  'class_to_idx': model.class_to_idx,
                  'state_dict': model.state_dict(),
                  'arch': args.arch,
                  'epochs': args.epochs,
                  'gpu': args.gpu,
                  'learning_rate': args.learning_rate,
                  'data_dir': args.data_dir}

    torch.save(checkpoint, args.checkpoint_dir)
    print('Model saved here: ', args.checkpoint_dir)


if __name__ == "__main__":
    args = parse_args()
    train_image_data, valid_image_data = load_data(args.data_dir)
    train_loader, valid_loader = get_loader(train_image_data, valid_image_data)
    model = create_model(args.arch, args.hidden_units)
    train(model, args.learning_rate, args.epochs, args.gpu)
    save_model(model, train_image_data, args)
```

▸ predict.py 2

▸ README.md

RETURN TO PATH

Student FAQ